

# Introduction: roadmap

1.1 what *is* the Internet?

1.2 network edge

- end systems, access networks, links

1.3 network core


- packet switching, circuit switching, network structure

1.4 delay, loss, throughput in networks

1.5 protocol layers, service model

# Sources of packet delay and loss?


# Packet Switching: queueing delay, loss



## queuing and loss:

- ❖ If arrival rate (in bits) to link exceeds transmission rate of link for a period of time:
  - packets will queue, wait to be transmitted on link
  - packets can be dropped (lost) if memory (buffer) fills up
- ❖ Do you want to keep this buffer full, empty, or what?

# Four sources of packet delay



$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$


## $d_{\text{proc}}$ : nodal processing

- check bit errors
- determine output link
- typically < msec
- **Negligible?**

## $d_{\text{queue}}$ : queueing delay

- time waiting at output link for transmission
- depends on congestion level of router

# Four sources of packet delay



$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

$d_{\text{trans}}$ : transmission delay:

- $L$ : packet length (bits)
- $R$ : link bandwidth (bps)
- $d_{\text{trans}} = L/R$

$d_{\text{prop}}$ : propagation delay:

- $d$ : length of physical link
- $s$ : propagation speed in medium ( $\sim 2 \times 10^8$  m/sec)
- $d_{\text{prop}} = d/s$

$d_{\text{trans}}$  and  $d_{\text{prop}}$   
very different

muhammadshahzad — -zsh — 80x24

```
64 bytes from 92.223.20.41: icmp_seq=15 ttl=46 time=162.168 ms
64 bytes from 92.223.20.41: icmp_seq=16 ttl=46 time=198.682 ms
64 bytes from 92.223.20.41: icmp_seq=17 ttl=46 time=155.489 ms
64 bytes from 92.223.20.41: icmp_seq=18 ttl=46 time=162.114 ms
64 bytes from 92.223.20.41: icmp_seq=19 ttl=46 time=169.983 ms
64 bytes from 92.223.20.41: icmp_seq=20 ttl=46 time=166.867 ms
64 bytes from 92.223.20.41: icmp_seq=21 ttl=46 time=168.033 ms
64 bytes from 92.223.20.41: icmp_seq=22 ttl=46 time=189.507 ms
64 bytes from 92.223.20.41: icmp_seq=23 ttl=46 time=184.747 ms
64 bytes from 92.223.20.41: icmp_seq=24 ttl=46 time=169.925 ms
64 bytes from 92.223.20.41: icmp_seq=25 ttl=46 time=174.059 ms
64 bytes from 92.223.20.41: icmp_seq=26 ttl=46 time=158.222 ms
64 bytes from 92.223.20.41: icmp_seq=27 ttl=46 time=166.997 ms
64 bytes from 92.223.20.41: icmp_seq=28 ttl=46 time=162.546 ms
64 bytes from 92.223.20.41: icmp_seq=29 ttl=46 time=169.475 ms
64 bytes from 92.223.20.41: icmp_seq=30 ttl=46 time=161.843 ms
64 bytes from 92.223.20.41: icmp_seq=31 ttl=46 time=174.667 ms
64 bytes from 92.223.20.41: icmp_seq=32 ttl=46 time=183.142 ms
64 bytes from 92.223.20.41: icmp_seq=33 ttl=46 time=194.294 ms
64 bytes from 92.223.20.41: icmp_seq=34 ttl=46 time=164.461 ms
64 bytes from 92.223.20.41: icmp_seq=35 ttl=46 time=155.234 ms
64 bytes from 92.223.20.41: icmp_seq=36 ttl=46 time=165.632 ms
64 bytes from 92.223.20.41: icmp_seq=37 ttl=46 time=171.418 ms
```

^C

muhammadshahzad — -zsh — 74x24

```
Request timeout for icmp_seq 2
64 bytes from 209.142.68.29: icmp_seq=3 ttl=118 time=81.927 ms
64 bytes from 209.142.68.29: icmp_seq=4 ttl=118 time=41.956 ms
64 bytes from 209.142.68.29: icmp_seq=5 ttl=118 time=60.169 ms
64 bytes from 209.142.68.29: icmp_seq=6 ttl=118 time=49.446 ms
64 bytes from 209.142.68.29: icmp_seq=7 ttl=118 time=66.815 ms
64 bytes from 209.142.68.29: icmp_seq=8 ttl=118 time=54.070 ms
64 bytes from 209.142.68.29: icmp_seq=9 ttl=118 time=66.603 ms
64 bytes from 209.142.68.29: icmp_seq=10 ttl=118 time=55.843 ms
64 bytes from 209.142.68.29: icmp_seq=11 ttl=118 time=54.960 ms
64 bytes from 209.142.68.29: icmp_seq=12 ttl=118 time=62.267 ms
64 bytes from 209.142.68.29: icmp_seq=13 ttl=118 time=138.473 ms
64 bytes from 209.142.68.29: icmp_seq=14 ttl=118 time=53.609 ms
64 bytes from 209.142.68.29: icmp_seq=15 ttl=118 time=75.197 ms
64 bytes from 209.142.68.29: icmp_seq=16 ttl=118 time=59.527 ms
64 bytes from 209.142.68.29: icmp_seq=17 ttl=118 time=71.403 ms
64 bytes from 209.142.68.29: icmp_seq=18 ttl=118 time=54.427 ms
64 bytes from 209.142.68.29: icmp_seq=19 ttl=118 time=51.863 ms
64 bytes from 209.142.68.29: icmp_seq=20 ttl=118 time=71.081 ms
64 bytes from 209.142.68.29: icmp_seq=21 ttl=118 time=47.509 ms
64 bytes from 209.142.68.29: icmp_seq=22 ttl=118 time=37.661 ms
64 bytes from 209.142.68.29: icmp_seq=23 ttl=118 time=42.590 ms
64 bytes from 209.142.68.29: icmp_seq=24 ttl=118 time=42.442 ms
```


^C

*A game server in Europe*

*A server in Chicago*

# Queueing delay (revisited)

- ❖  $R$ : link bandwidth (bps)
- ❖  $L$ : packet length (bits)
- ❖  $a$ : average packet arrival rate



- ❖  $La/R \sim 0$ : avg. queueing delay = ?
- ❖  $La/R > 1$ : avg. queueing delay = ?
- ❖  $La/R \approx 1$ : avg. queueing delay = ? (**nature of traffic**)


$La/R \sim 0$



$La/R > 1$

# “Real” Internet delays and routes

- ❖ what do “real” Internet delay & loss look like?
- ❖ `traceroute` program: provides delay measurement from source to router along end-end Internet path towards destination. For all  $i$ :
  - sends three packets that will reach router  $i$  on path towards destination
  - router  $i$  will return packets to sender
  - sender times interval between transmission and reply.






# “Real” Internet delays, routes

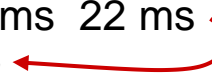
**traceroute:** gaia.cs.umass.edu to www.eurecom.fr

3 delay measurements from  
gaia.cs.umass.edu to cs-gw.cs.umass.edu



1 cs-gw (128.119.240.254) 1 ms 1 ms 2 ms  
2 border1-rt-fa5-1-0.gw.umass.edu (128.119.3.145) 1 ms 1 ms 2 ms  
3 cht-vbns.gw.umass.edu (128.119.3.130) 6 ms 5 ms 5 ms  
4 jn1-at1-0-0-19.wor.vbns.net (204.147.132.129) 16 ms 11 ms 13 ms  
5 jn1-so7-0-0-0.wae.vbns.net (204.147.136.136) 21 ms 18 ms 18 ms  
6 abilene-vbns.abilene.ucaid.edu (198.32.11.9) 22 ms 18 ms 22 ms  
7 nycm-wash.abilene.ucaid.edu (198.32.8.46) 22 ms 22 ms 22 ms  
8 62.40.103.253 (62.40.103.253) 104 ms 109 ms 106 ms  
9 de2-1.de1.de.geant.net (62.40.96.129) 109 ms 102 ms 104 ms  
10 de.fr1.fr.geant.net (62.40.96.50) 113 ms 121 ms 114 ms  
11 renater-gw.fr1.fr.geant.net (62.40.103.54) 112 ms 114 ms 112 ms  
12 nio-n2.cssi.renater.fr (193.51.206.13) 111 ms 114 ms 116 ms  
13 nice.cssi.renater.fr (195.220.98.102) 123 ms 125 ms 124 ms  
14 r3t2-nice.cssi.renater.fr (195.220.98.110) 126 ms 126 ms 124 ms  
15 eurecom-valbonne.r3t2.ft.net (193.48.50.54) 135 ms 128 ms 133 ms  
16 194.214.211.25 (194.214.211.25) 126 ms 128 ms 126 ms  
17 \* \* \*  
18 \* \* \*  
19 fantasia.eurecom.fr (193.55.113.142) 132 ms 128 ms 136 ms

trans-oceanic link




\* means no response (probe lost, router not replying)

\* Do some traceroutes from exotic countries at [www.traceroute.org](http://www.traceroute.org)


# Throughput

- ❖ *throughput*: rate (bits/time unit) at which bits transferred between sender/receiver
  - *instantaneous*: rate at given point in time
  - *average*: rate over longer period of time




# Throughput (more)

❖  $R_s < R_c$  What is average end-end throughput?



❖  $R_s > R_c$  What is average end-end throughput?




*bottleneck link*

link on end-end path that constrains end-end throughput

# Throughput: Internet scenario

- ❖ per-connection end-end throughput:  
 $\min(R_c, R_s, R/10)$
- ❖ in practice:  $R_c$  or  $R_s$  is often bottleneck



10 connections (fairly) share  
backbone bottleneck link  $R$  bits/sec

# Introduction: roadmap

1.1 what *is* the Internet?

1.2 network edge

- end systems, access networks, links

1.3 network core

- packet switching, circuit switching, network structure

1.4 delay, loss, throughput in networks

1.5 protocol layers, service models

# Protocol “layers”

*Networks are complex,  
with many “pieces”:*

- hosts
- routers
- links of various media
- applications
- protocols
- hardware, software


*Question:*

is there any hope of  
*organizing* structure of  
network?

.... or at least our  
discussion of networks?

# Internet protocol stack

- ❖ *application*: supporting network applications
  - FTP, SMTP, HTTP
- ❖ *transport*: process-process data transfer
  - TCP, UDP
- ❖ *network*: routing of datagrams from source to destination
  - IP, routing protocols
- ❖ *link*: data transfer between neighboring network elements
  - Ethernet, 802.11 (WiFi)
- ❖ *physical*: bits “on the wire”
- ❖ Which layers of protocol stack should be implemented by who?
- ❖ Which layers should be in HW and which in SW?




# Why layering?

dealing with complex systems:



- ❖ explicit structure allows identification, relationship of complex system's pieces
- ❖ modularization eases maintenance, updating of system
  - change of implementation of layer's service transparent to rest of system
- ❖ layering considered harmful?
  - Replication of functionality



# Encapsulation








*source*



switch

*destination*



router

# App-Layer: outline

## 2.1 principles of network applications

2.2 Web and HTTP

2.3 FTP

2.4 electronic mail

- SMTP, POP3, IMAP

2.5 DNS

2.6 P2P applications

2.7 socket programming  
with UDP and TCP

# Some network apps?

- ❖ e-mail
- ❖ web
- ❖ text messaging
- ❖ remote login
- ❖ P2P file sharing
- ❖ multi-user network games
- ❖ streaming stored video (YouTube, Hulu, Netflix)
- ❖ voice over IP (e.g., Skype)
- ❖ real-time video conferencing
- ❖ social networking
- ❖ search
- ❖ ...
- ❖ ...

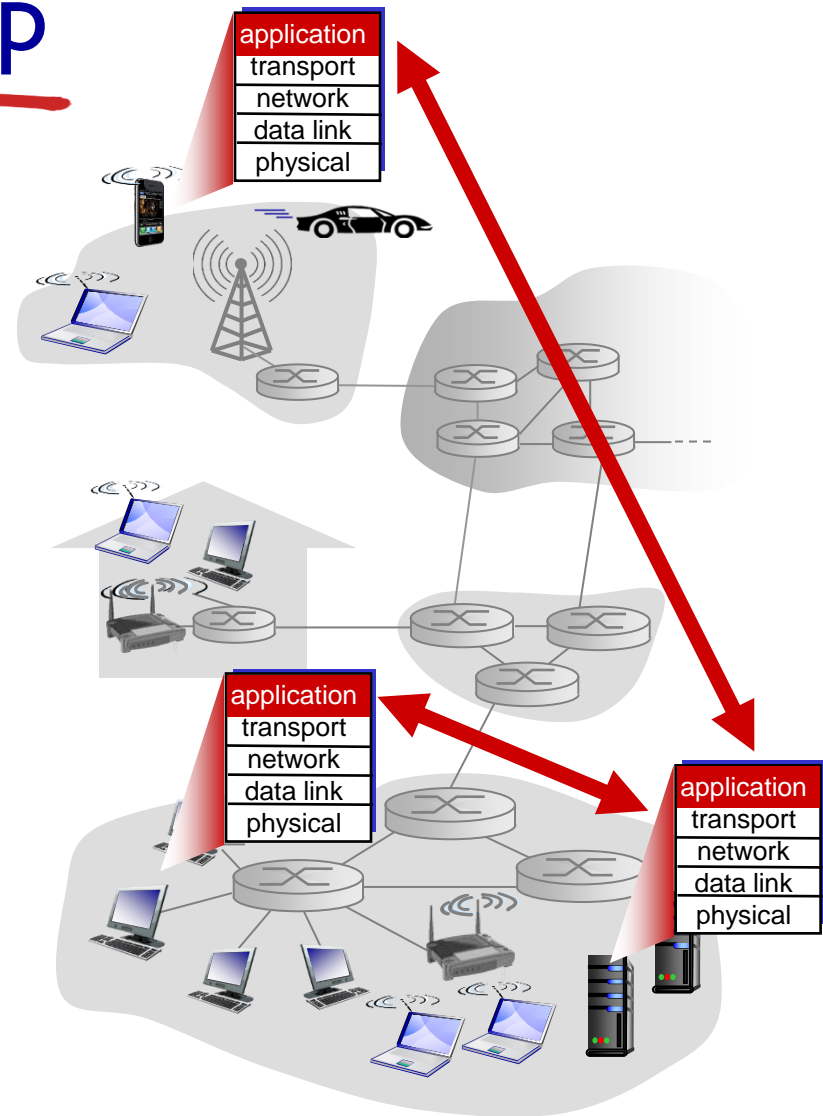
# Creating a network app

write programs that:

- ❖ run on (different) *end systems*
- ❖ communicate over network
- ❖ e.g., web server software communicates with browser software

no need to write software for  
network-core devices

- ❖ network-core devices do not run user applications
- ❖ applications on end systems allows for rapid app development, propagation




# Application architectures

possible structure of applications:

- ❖ client-server
- ❖ peer-to-peer (P2P)

# Client-server architecture



## server:


- ❖ always-on host
- ❖ permanent IP address
- ❖ data centers for scaling

## clients:

- ❖ communicate with server
- ❖ may be intermittently connected
- ❖ may have dynamic IP addresses
- ❖ do not communicate directly with each other

# P2P architecture

- ❖ no always-on server
- ❖ arbitrary end systems directly communicate
- ❖ peers request service from other peers, provide service in return to other peers
  - *self scalability* – new peers bring new service capacity, as well as new service demands
- ❖ peers are intermittently connected and change IP addresses
  - complex management



# Processes communicating

*process*: program running within a host

- ❖ within same host, two processes communicate using **inter-process communication** (defined by OS)
- ❖ processes in different hosts communicate by exchanging **messages** over a network

clients, servers

*client process*: process that initiates communication

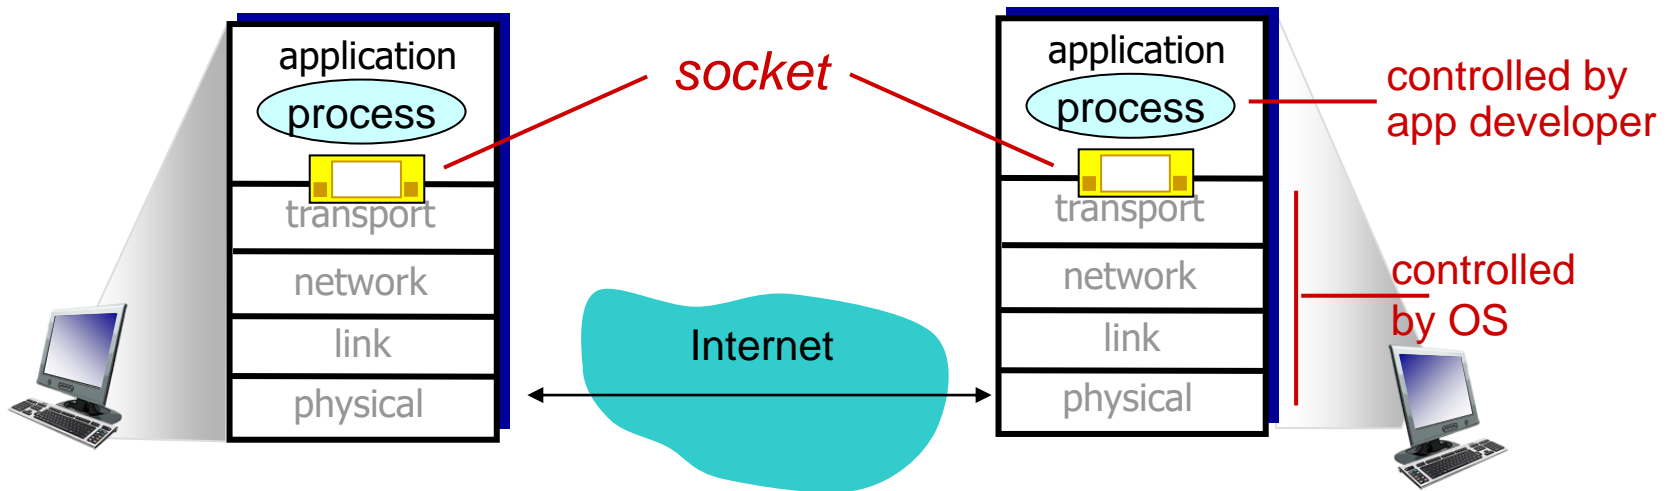
*server process*: process that waits to be contacted

- ❖ Applications with P2P architectures have client processes & server processes?



# Sockets

- ❖ process sends/receives messages to/from its **socket**
- ❖ socket analogous to door
  - sending process shoves message out door
  - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



# Addressing processes

- ❖ to receive messages, process must have *identifier*
- ❖ host device has unique 32-bit IP address
- ❖ Q: does IP address of host on which process runs suffice for identifying the process?
  - A: no, *many* processes can be running on same host
- ❖ *identifier* includes both **IP address** and **port numbers** associated with process on host.
- ❖ example port numbers:
  - HTTP server: 80
  - mail server: 25
- ❖ to send HTTP message to msu.edu web server:
  - **IP address**: 35.9.247.69
  - **port number**: 80

# App-layer protocol defines

- ❖ Types of messages exchanged,
  - e.g., request, response
- ❖ message syntax:
  - what fields in messages & how fields are delineated
- ❖ message semantics
  - meaning of information in fields
- ❖ rules for when and how processes send & respond to messages

## open protocols:

- ❖ defined in RFCs
- ❖ allows for interoperability
- ❖ e.g., HTTP, SMTP

## proprietary protocols:

- ❖ e.g., Skype
- ❖ Difference between Network application and application layer protocol?

# What transport service does an app need?

## data integrity

- ❖ some apps require 100% reliable data transfer; other apps can tolerate some loss
- ❖ Examples?
  - File transfer, web transactions; audio

## timing

- ❖ some apps require low delay to be “effective”; others don’t
- ❖ Examples?
  - Internet telephony, interactive games

## throughput

- ❖ some apps require minimum amount of throughput to be “effective”; others don’t
- ❖ Examples?
  - Multimedia

## security

- ❖ encryption, data integrity, ...

# Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	yes, 100' s msec
stored audio/video	loss-tolerant	same as above	yes
interactive games	loss-tolerant	few kbps up	yes, 100' s msec
text messaging	no loss	elastic	

# Internet transport protocols services

---

## TCP service:

- ❖ *reliable transport* between sending and receiving process
- ❖ *flow control*: sender won't overwhelm receiver
- ❖ *congestion control*: throttle sender when network overloaded
- ❖ *does not provide*: timing, minimum throughput guarantee, security
- ❖ *connection-oriented*: setup required between client and server processes

## UDP service:

- ❖ *unreliable data transfer* between sending and receiving process
- ❖ *does not provide*: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

Q: why bother? Why is there a UDP?

# Internet apps: application, transport protocols

	<b>application</b>	<b>application layer protocol</b>	<b>underlying transport protocol</b>
	e-mail	SMTP [RFC 2821]	TCP
remote terminal access		Telnet [RFC 854]	TCP
	Web	HTTP [RFC 2616]	TCP
	file transfer	FTP [RFC 959]	TCP
streaming multimedia		HTTP (e.g., YouTube), RTP [RFC 1889]	TCP or UDP
Internet telephony		SIP, RTP, proprietary (e.g., Skype)	TCP or UDP

# Securing TCP

## TCP & UDP

- ❖ no encryption
- ❖ cleartext passwds sent into socket traverse Internet in cleartext

## SSL

- ❖ provides encrypted TCP connection
- ❖ data integrity
- ❖ end-point authentication

## SSL is at app layer

- ❖ Apps use SSL libraries, which “talk” to TCP

## SSL socket API

- ❖ cleartext passwds sent into socket traverse Internet encrypted

Q: Why not run SSL over UDP?